

An Introduction to Apache, PySpark and Dataframe Transformations

A Comprehensive Guide to Master Big Data Analysis



Victor Roman

[Follow](#)

Jun 12, 2019 · 10 min read



Introduction: The Big Data Problem

Apache arises as a new engine and programming model for data analytics. Its origin goes back to 2009, and the main reasons why it has gained so much importance in the past recent years are due to changes in enconomic factors that underline computer applications and hardware.

Historically, the power of computers only grew with time. Each year, new processors were able to perform operations faster and the applications that run on top of them automatically got faster.

All of this changed in 2005, when the limits in heat disipation caused the switch from making individual processors faster, to start exploring the parallelization of CPU cores. This meant that applications and the code that run them must be changed too. All of this is what layed out the ground of new models like Apache Spark.

In addition, the cost of sensors and storing units only had decreased on the last years. Nowadays is completely unexpensive to collect and store vast amounts of information.

There is so much data available, that the way to process it and analyze it, must change radically too, by making large parallel computations on

clusters of computers. These clusters enable the synergic combination of those computers' power, simultaneously, and make much easier tackling expensive computational tasks like data processing.

And this is where Apache Spark comes into play.

What is Apache Spark

As found on the great book: Spark — The Definitive Guide:


“Apache Spark is a unified computing engine and a set of libraries for parallel data processing on clusters of computers”

Nowadays, Apache Spark is the most popular open source engine to Big Data processing. And the main reasons are:

- It supports programming languages as widely used as: Python, Scala, Java and R.
- It supports SQL tasks.
- It enables data streaming.

- It has libraries for Machine Learning and Deep Learning.
- It can be run in a single machine or in a cluster of computers.

The following is a sketch that illustrates the different libraries available in the Spark ecosystem.



How to Set Up and Run Apache Spark

Throughout these series of articles, we will focus on Apache Spark Python's library, PySpark. As stated before, Spark can be run both locally and in a cluster of computers. There are several ways to configure our machines to run Spark locally, but are out of the scope of these articles.

One of the simplest and fastest ways to work with PsyPark and unlock its immense processing power, is with the free website Databricks, concretely by using its Community Edition.

To get started we shoud simply go to:

Try Databricks

Unlimited clusters that can scale to any size Job scheduler to execute jobs for production pipelines Fully interactive...

databricks.com

And select its Community Edition:



Platform Solutions Customers Learn Partners Events Open Source Company



DATABRICKS PLATFORM – FREE TRIAL

For businesses looking for a zero-management cloud platform built around Apache Spark

COMMUNITY EDITION

For students and educational institutions just getting started with Apache Spark

- Unlimited clusters that can scale to any size
- Job scheduler to execute jobs for production pipelines
- Fully interactive notebook with collaboration, dashboards, REST APIs
- Advanced security, role-based access controls, and audit logs
- Single Sign On support
- Integration with BI tools such as Tableau, Qlik, and Looker
- 14-day full feature trial (excludes cloud charges)

GET STARTED


- Single cluster limited to 6GB and no worker nodes
- Basic notebook without collaboration
- Limited to 3 max users
- Public environment to share your work

GET STARTED

Then, we must create an account.

Running A Temporal Cluster

Once we have created an account, to be able to start working, we should create a temporary cluster.



The screenshot shows the Databricks Clusters interface. On the left is a sidebar with icons for Home, Workspace, Recents, Data, Clusters (which is selected), Jobs, and Search. The main area is titled 'Clusters' and contains a sub-section 'Interactive Clusters'. A table lists one cluster: 'Medium Sample' (Name), 'Running' (State), '1 (0 spot)' (Nodes), 'Community Optimized' (Driver), 'Community Optimized' (Worker), '5.3 (includes Apache Spar...' (Runtime), and 'v.roman.aragay@hotmail.com' (Creator). Below this is a section for 'Job Clusters' which says 'No clusters found'. At the top of the 'Interactive Clusters' section, there is a blue button labeled '+ Create Cluster' with a red box drawn around it to indicate it as the next step.

As it is a free version, these clusters have a default of 6 Gb of RAM and can be run for 6 hours each. In order to develop industrial projects or work with Data Pipelines, it is suggested to use the premium platform.


But for the aim of these tutorials, the community edition will be more than enough.


Adding Data

In order to add data to work with:

- Click on the data tab
- Then add data

You can work both with available data uploaded by other users or with data uploaded from your computer.





Once, done we can create a Table in a Notebook and we are all set up!


Pyspark Applications & Partitions

To understand how Apache Spark works we should talk about the core components of a Spark Application: The Driver, the Executors and the Cluster Manager.

The following is a very illustrative sketch of a Spark Application Architecture:

Driver Process

Executors



Driver

The driver is located in a node of the cluster of computers and performs three main tasks:

1. Holds information about the Spark Application
2. Responds to a input, for example a user's programm
3. Analyzes, distributes and schedules the tasks to be done by the executors.

Executors

The executors are the ones that actually perform the work assigned by the driver. They do two things:

1. Executing the code assigned to them.
2. Reporting the state of the computation to the driver.

Cluster Manager

The cluster manager is the responsible for:

1. Controlling the physical computers
2. Distributing resources to Spark Applications

There can be several Spark Applications running on the same cluster, at the same time, and all of them will be managed by the Cluster Manager.

PySpark Dataframes

Apache Spark works with several data abstractions, each with an specific interface to work with. The most common abstractions are:

- Datasets
- Dataframes

- SQL Tables
- Resilient Distributed Datasets

Throughout these series we will focus on the most common unit to represent and store data in Apache Spark, Dataframes.

Dataframes are data tables with rows and columns, the closest analogy to understand them are spreadsheets with labeled columns.

One important feature of Dataframes is their schema. A Dataframe's schema is a list with its columns names and the type of data that each column stores.

Other relevant attribute of Dataframes is that they are not located in one simple computer, in fact they can be splitted through hundreds of machines. This is due to optimize the processing of the information and when data is too large to fit a single machine.

Apache Partitions

As stated before, the executors perform the work assigned by the driver, and they do it in a parallel fashion, in order to be able to do this, Spark split data into different partitions.

These partitions are collections of rows located in a single computer within a cluster. When we talk about Dataframe's partitions we are talking about how the data is distributed across all the machines on our cluster.

Most of the time we will not specify explicitly how the partitions will be done in our clusters, but with our code we will transmit high-level transformations of the data and Spark will realize by itself which is the optimal way to perform these partitions. Always looking for obtaining the maximum processing efficiency.

Low level APIs to perform these operations are out of the scope of these series.

Dataframes Transformations

First of all, we have to understand that transformations are modifications that we specify to do to our dataframes.

These transformations are specified in a high-level fashion and will not be executed until we explicitly call for an action to be made.

This way of working is called lazy evaluation, and the aim is to improve efficiency. When we call for transformations to be made, Spark will design a plan to perform optimally these tasks, and will not execute it until the very last minute when we call an action (like `.show()` or `.collect()`)

Apple Stock Price

Now, we will explore some of the most common actions and transformations. We are going to work with Apple stock price's data, from 2010 to 2016. We will perform some exploratory data analysis, data transformations, deal with missing values and perform grouping and aggregating.

Import Dataframe

To initialize and display a dataframe, the code will be the following:

```
# File location and type
file_location = "/FileStore/tables/appl_stock.csv"
file_type = "csv"

# CSV options
infer_schema = "true"
first_row_is_header = "true"
delimiter = ","

# The applied options are for CSV files. For other file types, these
# will be ignored.
df = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)

# Display Dataframe
display(df)
```

```
▶ (3) Spark Jobs
▶ df: pyspark.sql.dataframe.DataFrame = [Date: timestamp, Open: double ... 5 more fields]
```

Date	Open	High	Low	Close	Volume	Adj Close
2010-01-04T00:00:00.000+0000	213.429998	214.499996	212.38000099999996	214.009998	123432400	27.727039
2010-01-05T00:00:00.000+0000	214.599998	215.589994	213.249994	214.379993	150476200	27.77497600000002
2010-01-06T00:00:00.000+0000	214.379993	215.23	210.750004	210.969995	138040000	27.33317800000004
2010-01-07T00:00:00.000+0000	211.75	212.000006	209.050005	210.58	119282800	27.28265
2010-01-08T00:00:00.000+0000	210.299994	212.000006	209.0600050000002	211.9800049999998	111902700	27.464034
2010-01-11T00:00:00.000+0000	212.7999970000002	213.000002	208.450005	210.1100029999998	115557400	27.221758
2010-01-12T00:00:00.000+0000	209.1899949999998	209.7699950000002	206.419998	207.720001	148614900	26.91211
2010-01-13T00:00:00.000+0000	207.870005	210.9299950000002	204.099998	210.650002	151473000	27.29172
2010-01-14T00:00:00.000+0000	210.4100000000002	210.4500007000002	206.020004	206.42	100000000	27.152657

Showing the first 1000 rows.



Get Dataframe's Schema

The schema of a dataframe is the description of the structure of the data, it is a collection of StructField objects and provides information about the type of the data in a dataframe.

To display the Dataframe's Schema is as simple as:

```
# Display Dataframe's Schema
df.printSchema()
```

```
root
|-- Date: timestamp (nullable = true)
|-- Open: double (nullable = true)
|-- High: double (nullable = true)
|-- Low: double (nullable = true)
|-- Close: double (nullable = true)
|-- Volume: integer (nullable = true)
|-- Adj Close: double (nullable = true)
```

Perform Filtering and Transformations

To filter our data, to get only those rows that have a closing price smaller than \$500, we could run the following line of code:

```
# Filter data usign pyspark
df.filter(" Close < 500").show()
```

Date	Open	High	Low	Close	Volume	Adj Close
2010-01-04 00:00:00	213.429998	214.499996	212.38000099999996	214.009998	123432400	27.727039
2010-01-05 00:00:00	214.599998	215.589994	213.249994	214.379993	150476200	27.774976000000002
2010-01-06 00:00:00	214.379993	215.23	210.750004	210.969995	138040000	27.333178000000004
2010-01-07 00:00:00	211.75	212.000006	209.050005	210.58	119282800	27.28265
2010-01-08 00:00:00	210.299994	212.000006	209.06000500000002	211.9800049999998	111902700	27.464034
2010-01-11 00:00:00	212.7999700000002	213.000002	208.450005	210.1100029999998	115557400	27.221758
2010-01-12 00:00:00	209.1899949999998	209.7699950000002	206.419998	207.720001	148614900	26.91211
2010-01-13 00:00:00	207.870005	210.9299950000002	204.099998	210.650002	151473000	27.29172
2010-01-14 00:00:00	210.1100029999998	210.4599970000002	209.020004	209.43	108223500	27.133657
2010-01-15 00:00:00	210.9299950000002	211.5999970000003	205.869999	205.93	148516900	26.68019799999997
2010-01-19 00:00:00	208.330002	215.1899990000003	207.240004	215.039995	182501900	27.86048499999997
2010-01-20 00:00:00	214.910006	215.549994	209.500002	211.73	153038200	27.431644
2010-01-21 00:00:00	212.079994	213.3099959999998	207.210003	208.069996	152038600	26.957455
2010-01-22 00:00:00	206.7800060000001	207.499996	197.16	197.75	220441900	25.620401
2010-01-25 00:00:00	202.5100020000001	204.699999	200.190002	203.070002	266424900	26.30965800000002
2010-01-26 00:00:00	205.9500010000001	213.710005	202.580004	205.940001	466777500	26.681494
2010-01-27 00:00:00	206.849995	210.58	199.530001	207.880005	430642100	26.93284000000002
2010-01-28 00:00:00	204.930004	205.500004	198.699995	199.289995	293375600	25.81992200000002
2010-01-29 00:00:00	201.079996	202.199995	190.250002	192.060003	311488100	24.883208
2010-02-01 00:00:00	192.3699969999998	196.0	191.2999989999999	194.729998	187469100	25.229131

We can also filter to only obtain certain columns:

```
# Filter data by columns
df.filter("Close < 500").select(['Open', 'Close']).show()
```

	Open	Close
	213.429998	214.009998
	214.599998	214.379993
	214.379993	210.969995
	211.75	210.58
	210.299994	211.98000499999998
	212.79999700000002	210.11000299999998
	209.18999499999998	207.720001
	207.870005	210.650002
	210.11000299999998	209.43
	210.92999500000002	205.93
	208.330002	215.039995
	214.910006	211.73
	212.079994	208.069996
	206.78000600000001	197.75
	202.51000200000001	203.070002
	205.95000100000001	205.940001
	206.849995	207.880005
	204.930004	199.289995
	201.079996	192.060003
	192.36999699999998	194.729998

only showing top 20 rows

To filter by one column and showing other, we will use the `.select()` method.

```
# Filter by one column and show other
df.filter(df['Close'] < 500).select('Volume').show()
```

Volume
123432400
123432400

```
| 150476200 |  
| 138040000 |  
| 119282800 |  
| 111902700 |  
| 115557400 |  
| 148614900 |  
| 151473000 |  
| 108223500 |  
| 148516900 |  
| 182501900 |  
| 153038200 |  
| 152038600 |  
| 220441900 |  
| 266424900 |  
| 466777500 |  
| 430642100 |  
| 293375600 |  
| 311488100 |  
| 187469100 |  
+-----+  
only showing top 20 rows
```

To filter by multiple conditions:

```
# Filter by multiple conditions: closing price < $200 and opening  
price > $200  
df.filter( (df['Close'] < 200) & (df['Open'] > 200) ).show()
```

Date	Open	High	Low	Close	Volume	Adj Close
2010-01-22 00:00:00	206.7800060000001	207.499996	197.16	197.75	220441900	25.620401
2010-01-28 00:00:00	204.930004	205.500004	198.699995	199.289995	293375600	25.819922000000002
2010-01-29 00:00:00	201.079996	202.199995	190.250002	192.060003	311488100	24.883208

Obtain a Statistic Summary of the Data

Similarly to other libraries like Pandas, we can obtain a statistic summary of the Dataframe by simply running the `.describe()` method.

```
# Display Statistic Summary  
df.describe().show()
```

summary	Open	High	Low	Close	Volume	Adj Close
count	1762	1762	1762	1762	1762	1762
mean	313.0763111589103	315.9112880164581	309.8282405079457	312.9270656379113	9.422577587968218E7	75.00174115607275
stddev	185.29946803981522	186.89817686485767	183.38391664371008	185.1471036170943	6.020518776592709E7	28.57492972179906
min	90.0	90.699997	89.470001	90.279999	11475900	24.881912
max	702.409988	705.070023	699.569977	702.100021	470249500	127.96609099999999

Add and Rename Columns

To add a new column to the dataframe, we will use the `.withColumn()` method as follows.

```
# Display Dataframe with new column  
df.withColumn('Doubled Adj Close', df['Adj Close']*2).select('Adj Close', 'Doubled Adj Close').show()
```

```

+-----+
|      27.727039|      55.454078|
|27.774976000000002|55.549952000000005|
|27.333178000000004| 54.66635600000001|
|      27.28265|      54.5653|
|      27.464034|      54.928068|
|      27.221758|      54.443516|
|      26.91211|      53.82422|
|      27.29172|      54.58344|
|      27.133657|      54.267314|
|26.680197999999997|53.360395999999994|
|27.860484999999997|55.720969999999994|
|      27.431644|      54.863288|
|      26.957455|      53.91491|
|      25.620401|      51.240802|
|26.309658000000002|52.619316000000005|
|      26.681494|      53.362988|
|26.932840000000002|53.865680000000005|
|25.819922000000002|51.639844000000004|
|      24.883208|      49.766416|
|      25.229131|      50.458262|
+-----+
only showing top 20 rows

```

To rename an existing column, we will use the `.withColumnRenamed()` method.

```

# Display Dataframe with renamed column
df.withColumnRenamed('Adj Close', 'Adjusted Close Price').show()

```

Date	Open	High	Low	Close	Volume	Adjusted Close Price
2010-01-04 00:00:00	213.429998	214.499996	212.3800009999996	214.009998	123432400	27.727039
2010-01-05 00:00:00	214.599998	215.589994	213.249994	214.379993	150476200	27.774976000000002
2010-01-06 00:00:00	214.379993	215.23	210.750004	210.969995	138040000	27.333178000000004
2010-01-07 00:00:00	211.75	212.000006	209.050005	210.58	119282800	27.28265

2010-01-08 00:00:00	210.299994	212.000006	209.06000500000002	211.9800499999998	111902700	27.464034
2010-01-11 00:00:00	212.7999970000002	213.000002	208.450005	210.1100029999998	115557400	27.221758
2010-01-12 00:00:00	209.1899949999998	209.7699950000002	206.419998	207.720001	148614900	26.912111
2010-01-13 00:00:00	207.870005	210.9299950000002	204.099998	210.650002	151473000	27.29172
2010-01-14 00:00:00	210.1100029999998	210.4599970000002	209.020004	209.43	108223500	27.133657
2010-01-15 00:00:00	210.9299950000002	211.5999970000003	205.869999	205.93	148516900	26.68019799999997
2010-01-19 00:00:00	208.330002	215.1899990000003	207.240004	215.039995	182501900	27.86048499999997
2010-01-20 00:00:00	214.910006	215.549994	209.500002	211.73	153038200	27.431644
2010-01-21 00:00:00	212.079994	213.3099959999998	207.210003	208.069996	152038600	26.957455
2010-01-22 00:00:00	206.7800060000001	207.499996	197.16	197.75	220441900	25.620401
2010-01-25 00:00:00	202.5100020000001	204.699999	200.190002	203.070002	266424900	26.30965800000002
2010-01-26 00:00:00	205.9500010000001	213.710005	202.580004	205.940001	466777500	26.681494
2010-01-27 00:00:00	206.849995	210.58	199.530001	207.880005	430642100	26.93284000000002
2010-01-28 00:00:00	204.930004	205.500004	198.699995	199.289995	293375600	25.81992200000002
2010-01-29 00:00:00	201.079996	202.199995	190.250002	192.060003	311488100	24.883208
2010-02-01 00:00:00	192.3699969999998	196.0	191.2999989999999	194.729998	187469100	25.229131

Grouping and Aggregating Data

Now, we will perform some grouping and aggregation of our data, in order to obtain meaningful insights. But first, we should import some libraries

```
# Import relevant libraries
from pyspark.sql.functions import
dayofmonth, hour, dayofyear, weekofyear, month, year, format_number, date_format,
rmat, mean, date_format, datediff, to_date, lit
```

Now, let us create a new column, with the year of each row:

```
# To know the average closing price per year
new_df = df.withColumn('Year', year(df['Date']))
new_df.show()
```

Date	Open	High	Low	Close	Volume	Adj Close	Year
2010-01-04 00:00:00	213.429998	214.499996	212.3800009999996	214.009998	123432400	27.727039	2010
2010-01-05 00:00:00	214.599998	215.589994	213.249994	214.379993	150476200	27.77497600000002	2010
2010-01-06 00:00:00	214.379993	215.23	210.750004	210.969995	138040000	27.33317800000004	2010
2010-01-07 00:00:00	211.75	212.000006	209.050005	210.58	119282800	27.28265	2010
2010-01-08 00:00:00	210.299994	212.000006	209.0600050000002	211.9800049999998	111902700	27.464034	2010
2010-01-11 00:00:00	212.7999970000002	213.000002	208.450005	210.1100029999998	115557400	27.221758	2010
2010-01-12 00:00:00	209.1899949999998	209.7699950000002	206.419998	207.720001	148614900	26.91211	2010
2010-01-13 00:00:00	207.870005	210.9299950000002	204.099998	210.650002	151473000	27.29172	2010
2010-01-14 00:00:00	210.1100029999998	210.4599970000002	209.020004	209.43	108223500	27.133657	2010
2010-01-15 00:00:00	210.9299950000002	211.5999970000003	205.869999	205.93	148516900	26.68019799999997	2010
2010-01-19 00:00:00	208.330002	215.1899990000003	207.240004	215.039995	182501900	27.86048499999997	2010
2010-01-20 00:00:00	214.910006	215.549994	209.500002	211.73	153038200	27.431644	2010
2010-01-21 00:00:00	212.079994	213.3099959999998	207.210003	208.069996	152038600	26.957455	2010
2010-01-22 00:00:00	206.7800060000001	207.499996	197.16	197.75	220441900	25.620401	2010
2010-01-25 00:00:00	202.5100020000001	204.699999	200.190002	203.070002	266424900	26.30965800000002	2010
2010-01-26 00:00:00	205.9500010000001	213.710005	202.580004	205.940001	466777500	26.681494	2010
2010-01-27 00:00:00	206.849995	210.58	199.530001	207.880005	430642100	26.93284000000002	2010
2010-01-28 00:00:00	204.930004	205.500004	198.699995	199.289995	293375600	25.81992200000002	2010
2010-01-29 00:00:00	201.079996	202.199995	190.250002	192.060003	311488100	24.883208	2010
2010-02-01 00:00:00	192.3699969999998	196.0	191.299989999999	194.729998	187469100	25.229131	2010

Now, lets group by this recently created 'Year' column and aggregate by the maximum, minimum and average prices of each year to obtain meaningful insights of the status and evolution of the price.

```
# Group and aggregate data
new_df.groupBy('Year').agg(f.max('Close').alias('Max Close'),
f.min('Close').alias('Min Close'), f.mean('Close').alias('Average Close')).orderBy('Year').show()
```

Year	Max Close	Min Close	Average Close
2010	325.470013	192.050003	259.842460000002
2011	422.2399980000007	315.320007	364.00432532142867
2012	702.100021	411.23	576.0497195640002
2013	570.090004	390.530006	472.6348802857143

```
| 2014 | 647.349983 | 90.279999 | 295.4023416507935 |
| 2015 | 133.0 | 103.120003 | 120.0399998055547 |
| 2016 | 118.25 | 90.339996 | 104.60400786904763 |
+-----+-----+-----+
```

We have achieved our goal! However, we still have some very difficult data to read. In fact we have way more decimals than we need.

Taking into account that we are working with prices of hundreds of dollars, more than two decimals do not provide us with relevant information.

So let's take advantage and learn to format the results to show us the number of decimals we want.

Formatting Our Data

To format our data we will use the `format_number()` function as follows:

```
# Import relevant functions
from pyspark.sql.functions import format_number, col

# Select the appropriate columns to format
cols = ['Max Close', 'Min Close', 'Average Close']

# Format the columns
formatted_df = new_df.select('Year', *[format_number(col(col_name), 2).name(col_name) for col_name in cols])
```

Year	Max Close	Min Close	Average Close
2010	325.47	192.05	259.84
2011	422.24	315.32	364.00
2012	702.10	411.23	576.05
2013	570.09	390.53	472.63
2014	647.35	90.28	295.40
2015	133.00	103.12	120.04
2016	118.25	90.34	104.60

User Defined Functions

Let's learn now how to apply functions defined by us to our dataframes. We will use it in this example to get a column with the month of the year in which each row was recorded.

```
# Import relevant functions
from pyspark.sql.functions import date_format, datediff, to_date,
lit, UserDefinedFunction, month
from pyspark.sql.types import StringType
from pyspark.sql import functions as F

# Create month list
month_lst = ['January', 'Feburary', 'March', 'April', 'May', 'June',
'July', 'August', 'September', 'October', 'November', 'December']

# Define the function
udf = UserDefinedFunction(lambda x: month_lst[int(x%12) - 1],
StringType())
```

```

# Add column to df with the number of the month of the year
df = df.withColumn('moy_number', month(df.Date))

# Apply function and generate a column with the name of the month of
# the year
df = df.withColumn('moy_name', udf("moy_number"))

```

	Date	Open	High	Low	Close	Volume	Adj Close	moy_number	moy_name
	2010-01-04 00:00:00	213.429998	214.499996	212.38000099999996	214.009998	123432400	27.727039	1	January
	2010-01-05 00:00:00	214.599998	215.589994	213.249994	214.379993	150476200	27.774976000000002	1	January
	2010-01-06 00:00:00	214.379993	215.23	210.750004	210.969995	138040000	27.333178000000004	1	January
	2010-01-07 00:00:00	211.75	212.000006	209.050005	210.58	119282800	27.28265	1	January
	2010-01-08 00:00:00	210.299994	212.000006	209.06000500000002	211.9800049999998	111902700	27.464034	1	January
	2010-01-11 00:00:00	212.7999970000002	213.000002	208.450005	210.1100029999998	115557400	27.221758	1	January
	2010-01-12 00:00:00	209.1899949999998	209.7699950000002	206.419998	207.720001	148614900	26.91211	1	January
	2010-01-13 00:00:00	207.870005	210.9299950000002	204.099998	210.650002	151473000	27.29172	1	January
	2010-01-14 00:00:00	210.1100029999998	210.4599970000002	209.020004	209.43	108223500	27.133657	1	January
	2010-01-15 00:00:00	210.9299950000002	211.5999970000003	205.869999	205.93	148516900	26.68019799999997	1	January
	2010-01-19 00:00:00	208.330002	215.1899990000003	207.240004	215.039995	182501900	27.86048499999997	1	January
	2010-01-20 00:00:00	214.910006	215.549994	209.500002	211.73	153038200	27.431644	1	January
	2010-01-21 00:00:00	212.079994	213.3099959999998	207.210003	208.069996	152038600	26.957455	1	January
	2010-01-22 00:00:00	206.7800060000001	207.499996	197.16	197.75	220441900	25.620401	1	January
	2010-01-25 00:00:00	202.5100020000001	204.699999	200.190002	203.070002	266424900	26.30965800000002	1	January
	2010-01-26 00:00:00	205.9500010000001	213.710005	202.580004	205.940001	466777500	26.681494	1	January
	2010-01-27 00:00:00	206.849995	210.58	199.530001	207.880005	430642100	26.93284000000002	1	January
	2010-01-28 00:00:00	204.930004	205.500004	198.699995	199.289995	293375600	25.81992200000002	1	January
	2010-01-29 00:00:00	201.079996	202.199995	190.250002	192.060003	311488100	24.883208	1	January
	2010-02-01 00:00:00	192.3699969999998	196.0	191.2999989999999	194.729998	187469100	25.229131	2	February

Success!

Conclusion

Throughout this article we have covered:

- The basis of Apache Spark
- We have gained an intuition of why it is important and how it operates

- Perform analysis operations with PySpark and Dataframes

On the next articles we will learn how to apply Machine Learning in PySpark and apply this knowledge to some projects. Stay tuned!

[Big Data](#) [Data Science](#) [Pyspark](#) [Data Analysis](#) [Spark](#)

Discover Medium

Welcome to a place where words matter. On Medium, smart voices and original ideas take center stage - with no ads in sight. Watch

Make Medium yours

Follow all the topics you care about, and we'll deliver the best stories for you to your homepage and inbox. Explore

Become a member

Get unlimited access to the best stories on Medium — and support writers while you're at it. Just \$5/month. Upgrade

[About](#) [Help](#) [Legal](#)